

Modeling

December 11, 2023

During our study of the data and research on the possible model solutions, there is one transformer model approach caught our eye. This transformer model approach was designed by Wijkhuizen, M., in the Kaggle competition (2023). Our project team decided to follow Wijkhuizen, M.'s approach to create a transformer model as one of the models to test for this project. Our goal with this approach is to get a better understanding of the transformer model since Wijkhuizen, M.'s approach is to build a transformer model from scratch and not fine-tune a base model.

```
[ ]: !pip install tensorflow-addons
```

Collecting tensorflow-addons

Downloading tensorflow_addons-0.23.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (611 kB)
611.8/611.8

kB 10.8 MB/s eta 0:00:00

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow-addons) (23.2)

Collecting typeguard<3.0.0,>=2.7 (from tensorflow-addons)

Downloading typeguard-2.13.3-py3-none-any.whl (17 kB)

Installing collected packages: typeguard, tensorflow-addons

Successfully installed tensorflow-addons-0.23.0 typeguard-2.13.3

```
[ ]: import numpy as np
import pandas as pd
import tensorflow as tf
import tensorflow_addons as tfa
import matplotlib.pyplot as plt
import matplotlib as mpl
import seaborn as sn

from tqdm.notebook import tqdm
from sklearn.model_selection import train_test_split, GroupShuffleSplit

import glob
import sys
import os
import math
import gc
import sys
```

```
import sklearn
import scipy
```

```
/usr/local/lib/python3.10/dist-  
packages/tensorflow-addons/utils/tfa_eol_msg.py:23: UserWarning:
```

TensorFlow Addons (TFA) has ended development and introduction of new features. TFA has entered a minimal maintenance and release mode until a planned end of life in May 2024.

Please modify downstream libraries to take dependencies from other repositories in our TensorFlow community (e.g. Keras, Keras-CV, and Keras-NLP).

For more information see: <https://github.com/tensorflow/addons/issues/2807>

```
warnings.warn(  

```

```
[ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: googledrive_dir = '/content/drive/MyDrive/Colab Notebooks/Data/asl-signs/'
```

```
[ ]: X_train = np.load(f'{googledrive_dir}/X_train.npy')  
y_train = np.load(f'{googledrive_dir}/y_train.npy')  
NON_EMPTY_FRAME_IDXS_TRAIN = np.load(f'{googledrive_dir}/  
↳NON_EMPTY_FRAME_IDXS_TRAIN.npy')  
X_val = np.load(f'{googledrive_dir}/X_val.npy')  
y_val = np.load(f'{googledrive_dir}/y_val.npy')  
NON_EMPTY_FRAME_IDXS_VAL = np.load(f'{googledrive_dir}/NON_EMPTY_FRAME_IDXS_VAL.  
↳npy')
```

```
[ ]: X_train.shape
```

```
[ ]: (80229, 64, 66, 3)
```

```
[ ]: y_train.shape
```

```
[ ]: (80229,)
```

```
[ ]: X_train.dtype
```

```
[ ]: dtype('float32')
```

```
[ ]: y_train.dtype
```

```
[ ]: dtype('int32')
```

```
[ ]: NON_EMPTY_FRAME_IDXS_TRAIN.shape
```

```
[ ]: (80229, 64)
```

```
[ ]: display(pd.Series(y_train).value_counts().to_frame('Class Count').  
↳iloc[[0,1,2,3,4, -5,-4,-3,-2,-1]])
```

	Class Count
135	358
136	352
60	351
194	351
67	348
170	277
249	267
56	266
21	263
231	255

```
[ ]: y_train
```

```
[ ]: array([ 25, 232, 48, ..., 86, 188, 105], dtype=int32)
```

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training  
↳  
# Epsilon value for layer normalisation  
LAYER_NORM_EPS = 1e-6  
  
# Dense layer units for landmarks  
LIPS_UNITS = 384  
HANDS_UNITS = 384  
POSE_UNITS = 384  
# final embedding and transformer embedding size  
UNITS = 512  
  
# Transformer  
NUM_BLOCKS = 2  
MLP_RATIO = 2  
  
# Dropout  
EMBEDDING_DROPOUT = 0.00  
MLP_DROPOUT_RATIO = 0.30  
CLASSIFIER_DROPOUT_RATIO = 0.10  
  
# Initializers  
INIT_HE_UNIFORM = tf.keras.initializers.he_uniform  
INIT_GLOROT_UNIFORM = tf.keras.initializers.glorot_uniform  
INIT_ZEROS = tf.keras.initializers.constant(0.0)  
# Activations  
GELU = tf.keras.activations.gelu
```

```
print(f'UNITS: {UNITS}')
```

UNITS: 512

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training
      ↪gislr-tf-data-processing-transformer-training
      # If True, processing data from scratch
      # If False, loads preprocessed data
      PREPROCESS_DATA = False
      TRAIN_MODEL = True
      # True: use 10% of participants as validation set
      # False: use all data for training -> gives better LB result
      USE_VAL = False

      N_ROWS = 543
      N_DIMS = 3
      DIM_NAMES = ['x', 'y', 'z']
      SEED = 42
      NUM_CLASSES = 250
      IS_INTERACTIVE = True
      VERBOSE = 1 if IS_INTERACTIVE else 2

      INPUT_SIZE = 64

      BATCH_ALL_SIGNS_N = 4
      BATCH_SIZE = 256
      N_EPOCHS = 100
      LR_MAX = 1e-3
      N_WARMUP_EPOCHS = 0
      WD_RATIO = 0.05
      MASK_VAL = 4237
```

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training
      ↪gislr-tf-data-processing-transformer-training
      USE_TYPES = ['left_hand', 'pose', 'right_hand']
      START_IDX = 468
      LIPS_IDXS0 = np.array([
          61, 185, 40, 39, 37, 0, 267, 269, 270, 409,
          291, 146, 91, 181, 84, 17, 314, 405, 321, 375,
          78, 191, 80, 81, 82, 13, 312, 311, 310, 415,
          95, 88, 178, 87, 14, 317, 402, 318, 324, 308,
      ])
      # Landmark indices in original data
      LEFT_HAND_IDXS0 = np.arange(468, 489)
      RIGHT_HAND_IDXS0 = np.arange(522, 543)
      LEFT_POSE_IDXS0 = np.array([502, 504, 506, 508, 510])
```

```

RIGHT_POSE_IDXS0 = np.array([503, 505, 507, 509, 511])
LANDMARK_IDXS_LEFT_DOMINANT0 = np.concatenate((LIPS_IDXS0, LEFT_HAND_IDXS0,
↳LEFT_POSE_IDXS0))
LANDMARK_IDXS_RIGHT_DOMINANT0 = np.concatenate((LIPS_IDXS0, RIGHT_HAND_IDXS0,
↳RIGHT_POSE_IDXS0))
HAND_IDXS0 = np.concatenate((LEFT_HAND_IDXS0, RIGHT_HAND_IDXS0), axis=0)
N_COLS = LANDMARK_IDXS_LEFT_DOMINANT0.size
# Landmark indices in processed data
LIPS_IDXS = np.argwhere(np.isin(LANDMARK_IDXS_LEFT_DOMINANT0, LIPS_IDXS0)).
↳squeeze()
LEFT_HAND_IDXS = np.argwhere(np.isin(LANDMARK_IDXS_LEFT_DOMINANT0,
↳LEFT_HAND_IDXS0)).squeeze()
RIGHT_HAND_IDXS = np.argwhere(np.isin(LANDMARK_IDXS_LEFT_DOMINANT0,
↳RIGHT_HAND_IDXS0)).squeeze()
HAND_IDXS = np.argwhere(np.isin(LANDMARK_IDXS_LEFT_DOMINANT0, HAND_IDXS0)).
↳squeeze()
POSE_IDXS = np.argwhere(np.isin(LANDMARK_IDXS_LEFT_DOMINANT0, LEFT_POSE_IDXS0)).
↳squeeze()

print(f'# HAND_IDXS: {len(HAND_IDXS)}, N_COLS: {N_COLS}')
```

HAND_IDXS: 21, N_COLS: 66

```

[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
LIPS_START = 0
LEFT_HAND_START = LIPS_IDXS.size
RIGHT_HAND_START = LEFT_HAND_START + LEFT_HAND_IDXS.size
POSE_START = RIGHT_HAND_START + RIGHT_HAND_IDXS.size

print(f'LIPS_START: {LIPS_START}, LEFT_HAND_START: {LEFT_HAND_START},
↳RIGHT_HAND_START: {RIGHT_HAND_START}, POSE_START: {POSE_START}')
```

LIPS_START: 0, LEFT_HAND_START: 40, RIGHT_HAND_START: 61, POSE_START: 61

```

[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
def get_lips_mean_std():
    # LIPS
    LIPS_MEAN_X = np.zeros([LIPS_IDXS.size], dtype=np.float32)
    LIPS_MEAN_Y = np.zeros([LIPS_IDXS.size], dtype=np.float32)
    LIPS_STD_X = np.zeros([LIPS_IDXS.size], dtype=np.float32)
    LIPS_STD_Y = np.zeros([LIPS_IDXS.size], dtype=np.float32)

    fig, axes = plt.subplots(3, 1, figsize=(15, N_DIMS*6))
```

```

    for col, ll in enumerate(tqdm( np.transpose(X_train[:, :, LIPS_IDXS],
↪[2,3,0,1]).reshape([LIPS_IDXS.size, N_DIMS, -1]) )):
        for dim, l in enumerate(ll):
            v = l[np.nonzero(l)]
            if dim == 0: # X
                LIPS_MEAN_X[col] = v.mean()
                LIPS_STD_X[col] = v.std()
            if dim == 1: # Y
                LIPS_MEAN_Y[col] = v.mean()
                LIPS_STD_Y[col] = v.std()

            axes[dim].boxplot(v, notch=False, showfliers=False,
↪positions=[col], whis=[5,95])

    for ax, dim_name in zip(axes, DIM_NAMES):
        ax.set_title(f'Lips {dim_name.upper()} Dimension', size=24)
        ax.tick_params(axis='x', labelsize=8)
        ax.grid(axis='y')

    plt.subplots_adjust(hspace=0.50)
    plt.show()


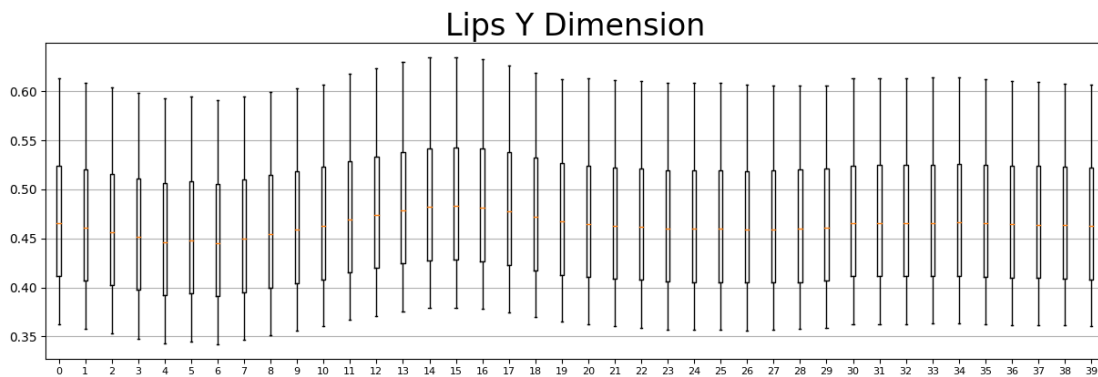
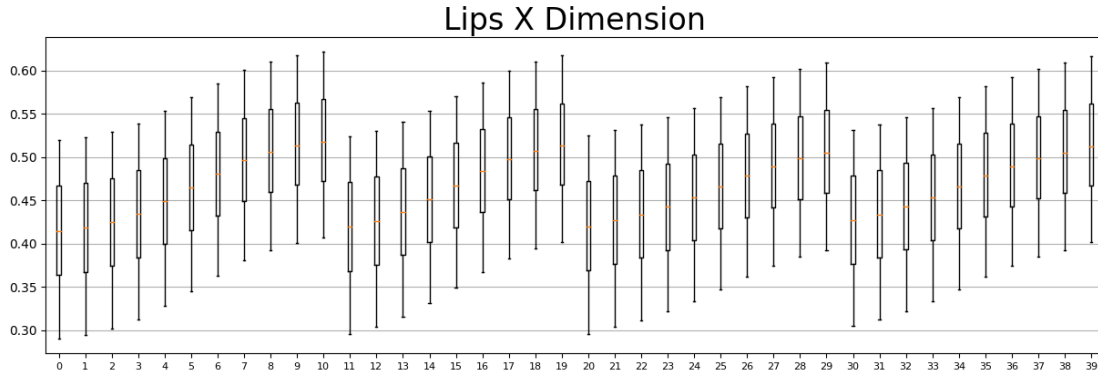
    LIPS_MEAN = np.array([LIPS_MEAN_X, LIPS_MEAN_Y]).T
    LIPS_STD = np.array([LIPS_STD_X, LIPS_STD_Y]).T

    return LIPS_MEAN, LIPS_STD

LIPS_MEAN, LIPS_STD = get_lips_mean_std()

```

0%| | 0/40 [00:00<?, ?it/s]



```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training
def get_left_right_hand_mean_std():
    # LEFT HAND
    LEFT_HANDS_MEAN_X = np.zeros([LEFT_HAND_IDXS.size], dtype=np.float32)
    LEFT_HANDS_MEAN_Y = np.zeros([LEFT_HAND_IDXS.size], dtype=np.float32)
    LEFT_HANDS_STD_X = np.zeros([LEFT_HAND_IDXS.size], dtype=np.float32)
```

```

LEFT_HANDS_STD_Y = np.zeros([LEFT_HAND_IDXS.size], dtype=np.float32)

fig, axes = plt.subplots(3, 1, figsize=(15, N_DIMS*6))

for col, ll in enumerate(tqdm( np.transpose(X_train[:, :, LEFT_HAND_IDXS],
↪[2,3,0,1]).reshape([LEFT_HAND_IDXS.size, N_DIMS, -1]) )):
    for dim, l in enumerate(ll):
        v = l[np.nonzero(l)]
        if dim == 0: # X
            LEFT_HANDS_MEAN_X[col] = v.mean()
            LEFT_HANDS_STD_X[col] = v.std()
        if dim == 1: # Y
            LEFT_HANDS_MEAN_Y[col] = v.mean()
            LEFT_HANDS_STD_Y[col] = v.std()
        # Plot
        axes[dim].boxplot(v, notch=False, showfliers=False,
↪positions=[col], whis=[5,95])

    for ax, dim_name in zip(axes, DIM_NAMES):
        ax.set_title(f'Hands {dim_name.upper()} Dimension', size=24)
        ax.tick_params(axis='x', labelsize=8)
        ax.grid(axis='y')

plt.subplots_adjust(hspace=0.50)
plt.show()



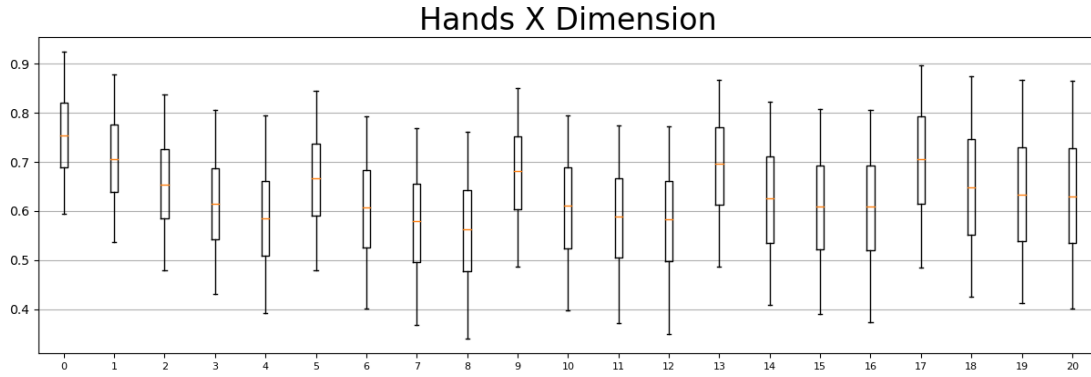
LEFT_HANDS_MEAN = np.array([LEFT_HANDS_MEAN_X, LEFT_HANDS_MEAN_Y]).T
LEFT_HANDS_STD = np.array([LEFT_HANDS_STD_X, LEFT_HANDS_STD_Y]).T

return LEFT_HANDS_MEAN, LEFT_HANDS_STD

LEFT_HANDS_MEAN, LEFT_HANDS_STD = get_left_right_hand_mean_std()

```

0%| | 0/21 [00:00<?, ?it/s]



```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training
def get_pose_mean_std():
    # POSE
    POSE_MEAN_X = np.zeros([POSE_IDXS.size], dtype=np.float32)
    POSE_MEAN_Y = np.zeros([POSE_IDXS.size], dtype=np.float32)
    POSE_STD_X = np.zeros([POSE_IDXS.size], dtype=np.float32)
```

```

POSE_STD_Y = np.zeros([POSE_IDXS.size], dtype=np.float32)

fig, axes = plt.subplots(3, 1, figsize=(15, N_DIMS*6))

for col, ll in enumerate(tqdm( np.transpose(X_train[:, :, POSE_IDXS],
↪ [2,3,0,1]).reshape([POSE_IDXS.size, N_DIMS, -1]) )):
    for dim, l in enumerate(ll):
        v = l[np.nonzero(l)]
        if dim == 0: # X
            POSE_MEAN_X[col] = v.mean()
            POSE_STD_X[col] = v.std()
        if dim == 1: # Y
            POSE_MEAN_Y[col] = v.mean()
            POSE_STD_Y[col] = v.std()

        axes[dim].boxplot(v, notch=False, showfliers=False,
↪ positions=[col], whis=[5,95])

    for ax, dim_name in zip(axes, DIM_NAMES):
        ax.set_title(f'Pose {dim_name.upper()} Dimension', size=24)
        ax.tick_params(axis='x', labelsize=8)
        ax.grid(axis='y')

plt.subplots_adjust(hspace=0.50)
plt.show()




POSE_MEAN = np.array([POSE_MEAN_X, POSE_MEAN_Y]).T
POSE_STD = np.array([POSE_STD_X, POSE_STD_Y]).T

return POSE_MEAN, POSE_STD

POSE_MEAN, POSE_STD = get_pose_mean_std()

```

0%| | 0/5 [00:00<?, ?it/s]



The attention mechanism is a key component in Transformer models, enabling the model to focus on different parts of the input sequence for each step of the output sequence. Attention enables the model to concentrate selectively on various segments of the input sequence for making predictions, rather than interpreting the entire sequence as a uniform-length vector. This feature has been crucial in the triumph of the transformer model, sparking extensive subsequent research and the development of numerous new models (Kumar, A. 2023). Wijkhuizen, M.'s custom transformer model deployed `attention_mask` in the Scaled Dot-Product function in a different way compared

to the classic transformer model in that this mask is applied in the Softmax step to selectively ignore or pay less attention to certain parts of the input, such as padding or irrelevant frames in a video sequence. Also, the Softmax layer was used instead of the Softmax function. The attention mechanism allows the model to focus on different parts of the input sequence dynamically, which is crucial for tasks like ASL recognition. In ASL, the importance of different landmarks can vary significantly across different signs. The multi-head attention mechanism is particularly well-suited to capture these varied dependencies.

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training
      ↪ https://stackoverflow.com/questions/67342988/verifying-the-implementation-of-multihead-attention-in-transformer
      # replaced softmax with softmax layer to support masked softmax
def scaled_dot_product(q,k,v, softmax, attention_mask):
    #calculates Q . K(transpose)
    qkt = tf.matmul(q,k,transpose_b=True)
    #caculates scaling factor
    dk = tf.math.sqrt(tf.cast(q.shape[-1],dtype=tf.float32))
    scaled_qkt = qkt/dk
    softmax = softmax(scaled_qkt, mask=attention_mask)

    z = tf.matmul(softmax,v)
    #shape: (m,Tx,depth), same shape as q,k,v
    return z

class MultiHeadAttention(tf.keras.layers.Layer):
    def __init__(self,d_model,num_of_heads):
        super(MultiHeadAttention,self).__init__()
        self.d_model = d_model
        self.num_of_heads = num_of_heads
        self.depth = d_model//num_of_heads
        self.wq = [tf.keras.layers.Dense(self.depth) for i in
↪range(num_of_heads)]
        self.wk = [tf.keras.layers.Dense(self.depth) for i in
↪range(num_of_heads)]
        self.wv = [tf.keras.layers.Dense(self.depth) for i in
↪range(num_of_heads)]
        self.wo = tf.keras.layers.Dense(d_model)
        self.softmax = tf.keras.layers.Softmax()

    def call(self,x, attention_mask):

        multi_attn = []
        for i in range(self.num_of_heads):
            Q = self.wq[i](x)
            K = self.wk[i](x)
            V = self.wv[i](x)
```

```

        multi_attn.append(scaled_dot_product(Q,K,V, self.softmax,
↪attention_mask))

        multi_head = tf.concat(multi_attn,axis=-1)
        multi_head_attention = self.wo(multi_head)
        return multi_head_attention

```

Wijkhuizen, M.'s transformer model approach is a custom-modified transformer model similar to the classic transformer model. There are custom changes to fit the ASL recognition tasks. This custom transformer only used a classic transformer encoder part for classification tasks, and it is the normal approach to only use the transformer encoder or decoder part alone. Besides the architecture part, there are two major customizations that Wijkhuizen, M. created that are different from the classic transformer model; one is the attention mechanism, and the other one is the embedding layer.

```

[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training
↪gislr-tf-data-processing-transformer-training
# Full Transformer
class Transformer(tf.keras.Model):
    def __init__(self, num_blocks):
        super(Transformer, self).__init__(name='transformer')
        self.num_blocks = num_blocks

    def build(self, input_shape):
        self.ln_1s = []
        self.mhas = []
        self.ln_2s = []
        self.mlps = []
        # Make Transformer Blocks
        for i in range(self.num_blocks):
            # Multi Head Attention
            self.mhas.append(MultiHeadAttention(UNITS, 8))
            # Multi Layer Perception
            self.mlps.append(tf.keras.Sequential([
                tf.keras.layers.Dense(UNITS * MLP_RATIO, activation=GELU,
↪kernel_initializer=INIT_GLOROT_UNIFORM),
                tf.keras.layers.Dropout(MLP_DROPOUT_RATIO),
                tf.keras.layers.Dense(UNITS,
↪kernel_initializer=INIT_HE_UNIFORM),
            ]))

    def call(self, x, attention_mask):
        # Iterate input over transformer blocks
        for mha, mlp in zip(self.mhas, self.mlps):
            x = x + mha(x, attention_mask)
            x = x + mlp(x)

```

```
return x
```

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
class LandmarkEmbedding(tf.keras.Model):
    def __init__(self, units, name):
        super(LandmarkEmbedding, self).__init__(name=f'{name}_embedding')
        self.units = units

    def build(self, input_shape):
        # Embedding for missing landmark in frame, initialized with zeros
        self.empty_embedding = self.add_weight(
            name=f'{self.name}_empty_embedding',
            shape=[self.units],
            initializer=INIT_ZEROS,
        )
        # Embedding
        self.dense = tf.keras.Sequential([
            tf.keras.layers.Dense(self.units, name=f'{self.name}_dense_1',
↳use_bias=False, kernel_initializer=INIT_GLOROT_UNIFORM),
            tf.keras.layers.Activation(GELU),
            tf.keras.layers.Dense(self.units, name=f'{self.name}_dense_2',
↳use_bias=False, kernel_initializer=INIT_HE_UNIFORM),
        ], name=f'{self.name}_dense')

    def call(self, x):
        return tf.where(
            # Checks whether landmark is missing in frame
            tf.reduce_sum(x, axis=2, keepdims=True) == 0,
            # If so, the empty embedding is used
            self.empty_embedding,
            # Otherwise the landmark data is embedded
            self.dense(x),
        )
```

Another customized part is the embedding layer. In the context of Transformer models, positional embeddings are crucial for providing information about the order or position of the elements in the input sequence (Huang, Z. et al., 2020). Compared to the classic transformer model, this ASL transformer model approach has an additional LandmarkEmbedding class that the Embedding class uses to deal with the embedding of individual landmarks. Each landmark type was embedded separately from lips, hand, and pose. The Embedding class is still handling the positional embedding.

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
class Embedding(tf.keras.Model):
    def __init__(self):
```

```

super(Embedding, self).__init__()

def get_diffs(self, l):
    S = l.shape[2]
    other = tf.expand_dims(l, 3)
    other = tf.repeat(other, S, axis=3)
    other = tf.transpose(other, [0,1,3,2])
    diffs = tf.expand_dims(l, 3) - other
    diffs = tf.reshape(diffs, [-1, INPUT_SIZE, S*S])
    return diffs

def build(self, input_shape):
    # Positional Embedding, initialized with zeros
    self.positional_embedding = tf.keras.layers.Embedding(INPUT_SIZE+1,
↳UNITS, embeddings_initializer=INIT_ZEROS)
    # Embedding layer for Landmarks
    self.lips_embedding = LandmarkEmbedding(LIPS_UNITS, 'lips')
    self.left_hand_embedding = LandmarkEmbedding(HANDS_UNITS, 'left_hand')
    self.pose_embedding = LandmarkEmbedding(POSE_UNITS, 'pose')
    # Landmark Weights
    self.landmark_weights = tf.Variable(tf.zeros([3], dtype=tf.float32),
↳name='landmark_weights')
    # Fully Connected Layers for combined landmarks
    self.fc = tf.keras.Sequential([
        tf.keras.layers.Dense(UNITS, name='fully_connected_1',
↳use_bias=False, kernel_initializer=INIT_GLOROT_UNIFORM),
        tf.keras.layers.Activation(GELU),
        tf.keras.layers.Dense(UNITS, name='fully_connected_2',
↳use_bias=False, kernel_initializer=INIT_HE_UNIFORM),
    ], name='fc')

    def call(self, lips0, left_hand0, pose0, non_empty_frame_idxs,
↳training=False):
        # Lips
        lips_embedding = self.lips_embedding(lips0)
        # Left Hand
        left_hand_embedding = self.left_hand_embedding(left_hand0)
        # Pose
        pose_embedding = self.pose_embedding(pose0)
        # Merge Embeddings of all landmarks with mean pooling
        x = tf.stack((
            lips_embedding, left_hand_embedding, pose_embedding,
        ), axis=3)
        x = x * tf.nn.softmax(self.landmark_weights)
        x = tf.reduce_sum(x, axis=3)
        # Fully Connected Layers

```

```

x = self.fc(x)
# Add Positional Embedding
max_frame_idx = tf.clip_by_value(
    tf.reduce_max(non_empty_frame_idx, axis=1, keepdims=True),
    1,
    np.PINF,
)
normalised_non_empty_frame_idx = tf.where(
    tf.math.equal(non_empty_frame_idx, -1.0),
    INPUT_SIZE,
    tf.cast(
        non_empty_frame_idx / max_frame_idx * INPUT_SIZE,
        tf.int32,
    ),
)
x = x + self.positional_embedding(normalised_non_empty_frame_idx)

return x

```

```

[ ]: # source:: https://stackoverflow.com/questions/60689185/
↳label-smoothing-for-sparse-categorical-crossentropy
def scce_with_ls(y_true, y_pred):
    # One Hot Encode Sparsely Encoded Target Sign
    y_true = tf.cast(y_true, tf.int32)
    y_true = tf.one_hot(y_true, NUM_CLASSES, axis=1)
    y_true = tf.squeeze(y_true, axis=2)
    # Categorical Crossentropy with native label smoothing support
    return tf.keras.losses.categorical_crossentropy(y_true, y_pred,
↳label_smoothing=0.25)

```

```

[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
def get_model():
    # Inputs
    frames = tf.keras.layers.Input([INPUT_SIZE, N_COLS, N_DIMS], dtype=tf.
↳float32, name='frames')
    non_empty_frame_idx = tf.keras.layers.Input([INPUT_SIZE], dtype=tf.
↳float32, name='non_empty_frame_idx')
    # Padding Mask
    mask0 = tf.cast(tf.math.not_equal(non_empty_frame_idx, -1), tf.float32)
    mask0 = tf.expand_dims(mask0, axis=2)
    # Random Frame Masking
    mask = tf.where(
        (tf.random.uniform(tf.shape(mask0)) > 0.25) & tf.math.not_equal(mask0,
↳0.0),
        1.0,
        0.0,
    )

```



```

)
# Correct Samples Which are all masked now...
mask = tf.where(
    tf.math.equal(tf.reduce_sum(mask, axis=[1,2], keepdims=True), 0.0),
    mask0,
    mask,
)

"""
    left_hand: 468:489
    pose: 489:522
    right_hand: 522:543
"""

x = frames
x = tf.slice(x, [0,0,0,0], [-1,INPUT_SIZE, N_COLS, 2])
# LIPS
lips = tf.slice(x, [0,0,LIPS_START,0], [-1,INPUT_SIZE, 40, 2])
lips = tf.where(
    tf.math.equal(lips, 0.0),
    0.0,
    (lips - LIPS_MEAN) / LIPS_STD,
)
# LEFT HAND
left_hand = tf.slice(x, [0,0,40,0], [-1,INPUT_SIZE, 21, 2])
left_hand = tf.where(
    tf.math.equal(left_hand, 0.0),
    0.0,
    (left_hand - LEFT_HANDS_MEAN) / LEFT_HANDS_STD,
)
# POSE
pose = tf.slice(x, [0,0,61,0], [-1,INPUT_SIZE, 5, 2])
pose = tf.where(
    tf.math.equal(pose, 0.0),
    0.0,
    (pose - POSE_MEAN) / POSE_STD,
)

# Flatten
lips = tf.reshape(lips, [-1, INPUT_SIZE, 40*2])
left_hand = tf.reshape(left_hand, [-1, INPUT_SIZE, 21*2])
pose = tf.reshape(pose, [-1, INPUT_SIZE, 5*2])

# Embedding
x = Embedding()(lips, left_hand, pose, non_empty_frame_idx)

# Encoder Transformer Blocks

```

```

x = Transformer(NUM_BLOCKS)(x, mask)

# Pooling
x = tf.reduce_sum(x * mask, axis=1) / tf.reduce_sum(mask, axis=1)
# Classifier Dropout
x = tf.keras.layers.Dropout(CLASSIFIER_DROPOUT_RATIO)(x)
# Classification Layer
x = tf.keras.layers.Dense(NUM_CLASSES, activation=tf.keras.activations.
↳softmax, kernel_initializer=INIT_GLOROT_UNIFORM)(x)

outputs = x

# Create Tensorflow Model
model = tf.keras.models.Model(inputs=[frames, non_empty_frame_idxs],
↳outputs=outputs)

# Sparse Categorical Cross Entropy With Label Smoothing
loss = scce_with_ls

# Adam Optimizer with weight decay
optimizer = tfa.optimizers.AdamW(learning_rate=1e-3, weight_decay=1e-5,
↳clipnorm=1.0)

# TopK Metrics
metrics = [
    tf.keras.metrics.SparseCategoricalAccuracy(name='acc'),
    tf.keras.metrics.SparseTopKCategoricalAccuracy(k=5, name='top_5_acc'),
    tf.keras.metrics.SparseTopKCategoricalAccuracy(k=10, name='top_10_acc'),
]

model.compile(loss=loss, optimizer=optimizer, metrics=metrics)

return model

```

```

[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
tf.keras.backend.clear_session()

model = get_model()

```

```

[ ]: # Plot model summary
model.summary(expand_nested=True)

```

Model: "model"

```

-----
Layer (type)                Output Shape              Param #   Connected to

```

```

=====
=====
non_empty_frame_idx (InputLayer) [(None, 64)] 0 []
tLayer)

tf.math.not_equal (TFOpLambda) (None, 64) 0
['non_empty_frame_idx[0][0]']
bda)

tf.cast (TFOpLambda) (None, 64) 0
['tf.math.not_equal[0][0]']

tf.expand_dims (TFOpLambda) (None, 64, 1) 0
['tf.cast[0][0]']
)

frames (InputLayer) [(None, 64, 66, 3)] 0 []

tf.compat.v1.shape (TFOpLambda) (3,) 0
['tf.expand_dims[0][0]']
mbda)

tf.slice (TFOpLambda) (None, 64, 66, 2) 0
['frames[0][0]']

tf.random.uniform (TFOpLambda) (None, 64, 1) 0
['tf.compat.v1.shape[0][0]']
bda)

tf.slice_1 (TFOpLambda) (None, 64, 40, 2) 0
['tf.slice[0][0]']

tf.slice_2 (TFOpLambda) (None, 64, 21, 2) 0
['tf.slice[0][0]']

tf.slice_3 (TFOpLambda) (None, 64, 5, 2) 0
['tf.slice[0][0]']

tf.math.greater (TFOpLambda) (None, 64, 1) 0
['tf.random.uniform[0][0]']
a)

tf.math.not_equal_1 (TFOpLambda) (None, 64, 1) 0
['tf.expand_dims[0][0]']
ambda)

tf.math.subtract (TFOpLambda) (None, 64, 40, 2) 0
['tf.slice_1[0][0]']

```

```

da)

tf.math.subtract_1 (TFOpLa (None, 64, 21, 2) 0
['tf.slice_2[0][0]']
mbda)

tf.math.subtract_2 (TFOpLa (None, 64, 5, 2) 0
['tf.slice_3[0][0]']
mbda)

tf.math.logical_and (TFOpL (None, 64, 1) 0
['tf.math.greater[0][0]',
ambda)
'tf.math.not_equal_1[0][0]']

tf.math.equal_1 (TFOpLambd (None, 64, 40, 2) 0
['tf.slice_1[0][0]']
a)

tf.math.truediv (TFOpLambd (None, 64, 40, 2) 0
['tf.math.subtract[0][0]']
a)

tf.math.equal_2 (TFOpLambd (None, 64, 21, 2) 0
['tf.slice_2[0][0]']
a)

tf.math.truediv_1 (TFOpLam (None, 64, 21, 2) 0
['tf.math.subtract_1[0][0]']
bda)

tf.math.equal_3 (TFOpLambd (None, 64, 5, 2) 0
['tf.slice_3[0][0]']
a)

tf.math.truediv_2 (TFOpLam (None, 64, 5, 2) 0
['tf.math.subtract_2[0][0]']
bda)

tf.where (TFOpLambda) (None, 64, 1) 0
['tf.math.logical_and[0][0]']

tf.where_2 (TFOpLambda) (None, 64, 40, 2) 0
['tf.math.equal_1[0][0]',
'tf.math.truediv[0][0]']

tf.where_3 (TFOpLambda) (None, 64, 21, 2) 0
['tf.math.equal_2[0][0]',

```

```

'tf.math.truediv_1[0][0] '

tf.where_4 (TFOpLambda)      (None, 64, 5, 2)      0
['tf.math.equal_3[0][0] ',
'tf.math.truediv_2[0][0] '

tf.math.reduce_sum (TFOpLa  (None, 1, 1)      0
['tf.where[0][0] '
mbda)

tf.reshape (TFOpLambda)      (None, 64, 80)      0
['tf.where_2[0][0] '

tf.reshape_1 (TFOpLambda)    (None, 64, 42)      0
['tf.where_3[0][0] '

tf.reshape_2 (TFOpLambda)    (None, 64, 10)      0
['tf.where_4[0][0] '

tf.math.equal (TFOpLambda)   (None, 1, 1)      0
['tf.math.reduce_sum[0][0] '

embedding (Embedding)        (None, 64, 512)      986243
['tf.reshape[0][0] ',
'tf.reshape_1[0][0] ',
'tf.reshape_2[0][0] ',
'non_empty_frame_idx[0][0] '
|-----|
|-----|
| embedding (Embedding)      multiple      33280      []
|
|
|
| lips_embedding (LandmarkE  multiple      178560      []
| mbedding)
|
|-----|
||-----|
||-----|
|| lips_embedding_dense (Se  (None, 64, 384)      178176      []
||
|| sequential)
||
||-----|
||-----|
||-----|
|| lips_embedding_dense_1    (None, 64, 384)      30720      []
||
|| (Dense)

```

```

|||
|||
|||
||| activation_1 (Activatio (None, 64, 384)          0          []
|||
||| n)
|||
|||
|||
||| lips_embedding_dense_2 (None, 64, 384)          147456      []
|||
||| (Dense)
|||
|||-----
||-----||
|-----
|-----
|-----|
| left_hand_embedding (Land multiple          163968      []
|
| markEmbedding)
|
||-----
||-----||
|| left_hand_embedding_dens (None, 64, 384)          163584      []
||
|| e (Sequential)
||
|||-----
|||-----|||
||| left_hand_embedding_den (None, 64, 384)          16128      []
|||
||| se_1 (Dense)
|||
|||
|||
||| activation_2 (Activatio (None, 64, 384)          0          []
|||
||| n)
|||
|||
|||
||| left_hand_embedding_den (None, 64, 384)          147456      []
|||
||| se_2 (Dense)
|||
|||-----
||-----||
|-----

```

```

-----|
| pose_embedding (LandmarkE multiple 151680 []
|
| mbedding)
|
|-----|
|-----||
|| pose_embedding_dense (Se (None, 64, 384) 151296 []
||
|| quential)
||
|-----|
|-----|||
||| pose_embedding_dense_1 (None, 64, 384) 3840 []
|||
||| (Dense)
|||
|||
|||
|||
||| activation_3 (Activatio (None, 64, 384) 0 []
|||
||| n)
|||
|||
|||
||| pose_embedding_dense_2 (None, 64, 384) 147456 []
|||
||| (Dense)
|||
|-----|
|-----||
|-----|
|-----|
| fc (Sequential) (None, 64, 512) 458752 []
|
|-----|
|-----||
|| fully_connected_1 (Dense (None, 64, 512) 196608 []
||
|| )
||
||
||
||
|| activation (Activation) (None, 64, 512) 0 []
||
||
||
|| fully_connected_2 (Dense (None, 64, 512) 262144 []

```

```

||
|| )
||
|-----|
|-----|
|-----|

tf.where_1 (TFOpLambda)      (None, 64, 1)      0
['tf.math.equal[0][0]',
'tf.expand_dims[0][0]',
'tf.where[0][0]']

transformer (Transformer)    (None, 64, 512)      4201472
['embedding[0][0]',
'tf.where_1[0][0]']
|-----|
|-----|
| multi_head_attention (Mul   multiple           1050624   []
|   tiHeadAttention)
|
|
|
| multi_head_attention_1 (M   multiple           1050624   []
|   ultiHeadAttention)
|
|
|
| sequential (Sequential)    (None, 64, 512)      1050112   []
|-----|
||-----||
|| dense_25 (Dense)          (None, 64, 1024)     525312    []
||
||
||
|| dropout (Dropout)         (None, 64, 1024)     0          []
||
||
||
|| dense_26 (Dense)          (None, 64, 512)      524800    []
||
|-----|
|-----|
| sequential_1 (Sequential)  (None, 64, 512)      1050112   []
|
|-----|
||-----||

```



```

-----||
|| dense_52 (Dense)          (None, 64, 1024)          525312    []
||
||
||
|| dropout_1 (Dropout)      (None, 64, 1024)          0          []
||
||
||
|| dense_53 (Dense)          (None, 64, 512)          524800    []
||
|-----|
-----|
-----

tf.math.multiply (TFOpLamb (None, 64, 512)          0
['transformer[0][0]',
da)
'tf.where_1[0][0]']

tf.math.reduce_sum_1 (TFOp (None, 512)          0
['tf.math.multiply[0][0]']
Lambda)

tf.math.reduce_sum_2 (TFOp (None, 1)          0
['tf.where_1[0][0]']
Lambda)

tf.math.truediv_3 (TFOpLam (None, 512)          0
['tf.math.reduce_sum_1[0][0]'],
bda)
'tf.math.reduce_sum_2[0][0]']

dropout (Dropout)          (None, 512)          0
['tf.math.truediv_3[0][0]']

dense (Dense)              (None, 250)          128250
['dropout[0][0]']

=====
=====
Total params: 5315965 (20.28 MB)
Trainable params: 5315965 (20.28 MB)
Non-trainable params: 0 (0.00 Byte)
-----
-----

```

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
def get_train_batch_all_signs(X, y, NON_EMPTY_FRAME_IDXS, n=BATCH_ALL_SIGNS_N):
    # Arrays to store batch in
    X_batch = np.zeros([NUM_CLASSES*n, INPUT_SIZE, N_COLS, N_DIMS], dtype=np.
↳float32)
    y_batch = np.arange(0, NUM_CLASSES, step=1/n, dtype=np.float32).astype(np.
↳int64)
    non_empty_frame_idxes_batch = np.zeros([NUM_CLASSES*n, INPUT_SIZE], dtype=np.
↳float32)

    # Dictionary mapping ordinally encoded sign to corresponding sample indices
    CLASS2IDX = {}
    for i in range(NUM_CLASSES):
        CLASS2IDX[i] = np.argwhere(y == i).squeeze().astype(np.int32)

    while True:
        # Fill batch arrays
        for i in range(NUM_CLASSES):
            idxs = np.random.choice(CLASS2IDX[i], n)
            X_batch[i*n:(i+1)*n] = X[idxs]
            non_empty_frame_idxes_batch[i*n:(i+1)*n] = NON_EMPTY_FRAME_IDXS[idxs]

        yield { 'frames': X_batch, 'non_empty_frame_idxes':
↳non_empty_frame_idxes_batch }, y_batch
```

```
[ ]: # Code From https://www.kaggle.com/code/markwijkhuizen/
↳gislr-tf-data-processing-transformer-training
# Actual Training
history = model.fit(
    x=get_train_batch_all_signs(X_train, y_train,
↳NON_EMPTY_FRAME_IDXS_TRAIN),
    steps_per_epoch=len(X_train) // (NUM_CLASSES * BATCH_ALL_SIGNS_N),
    epochs=N_EPOCHS,
    batch_size=BATCH_SIZE,
)
```

```
Epoch 1/100
80/80 [=====] - 34s 213ms/step - loss: 4.3454 - acc:
0.2691 - top_5_acc: 0.5258 - top_10_acc: 0.6316
Epoch 2/100
80/80 [=====] - 17s 212ms/step - loss: 3.3600 - acc:
0.5616 - top_5_acc: 0.8160 - top_10_acc: 0.8755
Epoch 3/100
80/80 [=====] - 17s 213ms/step - loss: 3.0562 - acc:
0.6665 - top_5_acc: 0.8765 - top_10_acc: 0.9160
Epoch 4/100
```

80/80 [=====] - 17s 212ms/step - loss: 2.9035 - acc: 0.7227 - top_5_acc: 0.9030 - top_10_acc: 0.9325
Epoch 5/100
80/80 [=====] - 17s 212ms/step - loss: 2.8028 - acc: 0.7592 - top_5_acc: 0.9164 - top_10_acc: 0.9422
Epoch 6/100
80/80 [=====] - 17s 212ms/step - loss: 2.7248 - acc: 0.7879 - top_5_acc: 0.9294 - top_10_acc: 0.9517
Epoch 7/100
80/80 [=====] - 17s 212ms/step - loss: 2.6656 - acc: 0.8100 - top_5_acc: 0.9388 - top_10_acc: 0.9578
Epoch 8/100
80/80 [=====] - 17s 213ms/step - loss: 2.6164 - acc: 0.8279 - top_5_acc: 0.9471 - top_10_acc: 0.9636
Epoch 9/100
80/80 [=====] - 17s 212ms/step - loss: 2.5814 - acc: 0.8413 - top_5_acc: 0.9530 - top_10_acc: 0.9684
Epoch 10/100
80/80 [=====] - 17s 213ms/step - loss: 2.5529 - acc: 0.8528 - top_5_acc: 0.9582 - top_10_acc: 0.9721
Epoch 11/100
80/80 [=====] - 17s 212ms/step - loss: 2.5209 - acc: 0.8652 - top_5_acc: 0.9637 - top_10_acc: 0.9753
Epoch 12/100
80/80 [=====] - 17s 212ms/step - loss: 2.4854 - acc: 0.8784 - top_5_acc: 0.9682 - top_10_acc: 0.9788
Epoch 13/100
80/80 [=====] - 17s 212ms/step - loss: 2.4557 - acc: 0.8899 - top_5_acc: 0.9715 - top_10_acc: 0.9814
Epoch 14/100
80/80 [=====] - 17s 212ms/step - loss: 2.4267 - acc: 0.9001 - top_5_acc: 0.9750 - top_10_acc: 0.9841
Epoch 15/100
80/80 [=====] - 17s 212ms/step - loss: 2.4016 - acc: 0.9093 - top_5_acc: 0.9784 - top_10_acc: 0.9862
Epoch 16/100
80/80 [=====] - 17s 213ms/step - loss: 2.3756 - acc: 0.9175 - top_5_acc: 0.9812 - top_10_acc: 0.9879
Epoch 17/100
80/80 [=====] - 17s 213ms/step - loss: 2.3534 - acc: 0.9265 - top_5_acc: 0.9839 - top_10_acc: 0.9898
Epoch 18/100
80/80 [=====] - 17s 212ms/step - loss: 2.3355 - acc: 0.9319 - top_5_acc: 0.9859 - top_10_acc: 0.9910
Epoch 19/100
80/80 [=====] - 17s 213ms/step - loss: 2.3125 - acc: 0.9420 - top_5_acc: 0.9882 - top_10_acc: 0.9927
Epoch 20/100

80/80 [=====] - 17s 212ms/step - loss: 2.3044 - acc: 0.9432 - top_5_acc: 0.9892 - top_10_acc: 0.9937
Epoch 21/100
80/80 [=====] - 17s 213ms/step - loss: 2.2896 - acc: 0.9491 - top_5_acc: 0.9905 - top_10_acc: 0.9948
Epoch 22/100
80/80 [=====] - 17s 213ms/step - loss: 2.2773 - acc: 0.9524 - top_5_acc: 0.9914 - top_10_acc: 0.9948
Epoch 23/100
80/80 [=====] - 17s 212ms/step - loss: 2.2823 - acc: 0.9523 - top_5_acc: 0.9921 - top_10_acc: 0.9953
Epoch 24/100
80/80 [=====] - 17s 213ms/step - loss: 2.2719 - acc: 0.9547 - top_5_acc: 0.9926 - top_10_acc: 0.9959
Epoch 25/100
80/80 [=====] - 17s 212ms/step - loss: 2.2757 - acc: 0.9541 - top_5_acc: 0.9927 - top_10_acc: 0.9959
Epoch 26/100
80/80 [=====] - 17s 213ms/step - loss: 2.2646 - acc: 0.9575 - top_5_acc: 0.9941 - top_10_acc: 0.9966
Epoch 27/100
80/80 [=====] - 17s 212ms/step - loss: 2.2591 - acc: 0.9590 - top_5_acc: 0.9944 - top_10_acc: 0.9970
Epoch 28/100
80/80 [=====] - 17s 212ms/step - loss: 2.2620 - acc: 0.9582 - top_5_acc: 0.9940 - top_10_acc: 0.9967
Epoch 29/100
80/80 [=====] - 17s 212ms/step - loss: 2.2553 - acc: 0.9603 - top_5_acc: 0.9952 - top_10_acc: 0.9973
Epoch 30/100
80/80 [=====] - 17s 212ms/step - loss: 2.2412 - acc: 0.9642 - top_5_acc: 0.9955 - top_10_acc: 0.9980
Epoch 31/100
80/80 [=====] - 17s 212ms/step - loss: 2.2417 - acc: 0.9638 - top_5_acc: 0.9961 - top_10_acc: 0.9982
Epoch 32/100
80/80 [=====] - 17s 212ms/step - loss: 2.2300 - acc: 0.9687 - top_5_acc: 0.9965 - top_10_acc: 0.9983
Epoch 33/100
80/80 [=====] - 17s 213ms/step - loss: 2.2245 - acc: 0.9690 - top_5_acc: 0.9968 - top_10_acc: 0.9985
Epoch 34/100
80/80 [=====] - 17s 213ms/step - loss: 2.2157 - acc: 0.9719 - top_5_acc: 0.9974 - top_10_acc: 0.9989
Epoch 35/100
80/80 [=====] - 17s 213ms/step - loss: 2.2132 - acc: 0.9727 - top_5_acc: 0.9979 - top_10_acc: 0.9990
Epoch 36/100

80/80 [=====] - 17s 212ms/step - loss: 2.2120 - acc: 0.9730 - top_5_acc: 0.9980 - top_10_acc: 0.9991
Epoch 37/100
80/80 [=====] - 17s 213ms/step - loss: 2.2113 - acc: 0.9730 - top_5_acc: 0.9976 - top_10_acc: 0.9991
Epoch 38/100
80/80 [=====] - 17s 212ms/step - loss: 2.1991 - acc: 0.9768 - top_5_acc: 0.9981 - top_10_acc: 0.9993
Epoch 39/100
80/80 [=====] - 17s 212ms/step - loss: 2.2091 - acc: 0.9746 - top_5_acc: 0.9982 - top_10_acc: 0.9993
Epoch 40/100
80/80 [=====] - 17s 213ms/step - loss: 2.2064 - acc: 0.9745 - top_5_acc: 0.9981 - top_10_acc: 0.9993
Epoch 41/100
80/80 [=====] - 17s 213ms/step - loss: 2.1985 - acc: 0.9768 - top_5_acc: 0.9986 - top_10_acc: 0.9995
Epoch 42/100
80/80 [=====] - 17s 212ms/step - loss: 2.1908 - acc: 0.9786 - top_5_acc: 0.9984 - top_10_acc: 0.9994
Epoch 43/100
80/80 [=====] - 17s 212ms/step - loss: 2.1864 - acc: 0.9797 - top_5_acc: 0.9989 - top_10_acc: 0.9996
Epoch 44/100
80/80 [=====] - 17s 212ms/step - loss: 2.1894 - acc: 0.9793 - top_5_acc: 0.9987 - top_10_acc: 0.9995
Epoch 45/100
80/80 [=====] - 17s 213ms/step - loss: 2.1851 - acc: 0.9804 - top_5_acc: 0.9987 - top_10_acc: 0.9995
Epoch 46/100
80/80 [=====] - 17s 212ms/step - loss: 2.1856 - acc: 0.9799 - top_5_acc: 0.9989 - top_10_acc: 0.9996
Epoch 47/100
80/80 [=====] - 17s 213ms/step - loss: 2.1821 - acc: 0.9809 - top_5_acc: 0.9989 - top_10_acc: 0.9996
Epoch 48/100
80/80 [=====] - 17s 212ms/step - loss: 2.1774 - acc: 0.9813 - top_5_acc: 0.9990 - top_10_acc: 0.9995
Epoch 49/100
80/80 [=====] - 17s 212ms/step - loss: 2.1757 - acc: 0.9830 - top_5_acc: 0.9990 - top_10_acc: 0.9996
Epoch 50/100
80/80 [=====] - 17s 213ms/step - loss: 2.1845 - acc: 0.9802 - top_5_acc: 0.9990 - top_10_acc: 0.9995
Epoch 51/100
80/80 [=====] - 17s 212ms/step - loss: 2.1786 - acc: 0.9818 - top_5_acc: 0.9993 - top_10_acc: 0.9998
Epoch 52/100

80/80 [=====] - 17s 212ms/step - loss: 2.1759 - acc: 0.9815 - top_5_acc: 0.9990 - top_10_acc: 0.9996
Epoch 53/100
80/80 [=====] - 17s 212ms/step - loss: 2.1649 - acc: 0.9839 - top_5_acc: 0.9991 - top_10_acc: 0.9996
Epoch 54/100
80/80 [=====] - 17s 212ms/step - loss: 2.1734 - acc: 0.9827 - top_5_acc: 0.9992 - top_10_acc: 0.9997
Epoch 55/100
80/80 [=====] - 17s 213ms/step - loss: 2.1685 - acc: 0.9835 - top_5_acc: 0.9991 - top_10_acc: 0.9996
Epoch 56/100
80/80 [=====] - 17s 212ms/step - loss: 2.1836 - acc: 0.9807 - top_5_acc: 0.9991 - top_10_acc: 0.9997
Epoch 57/100
80/80 [=====] - 17s 212ms/step - loss: 2.1731 - acc: 0.9830 - top_5_acc: 0.9992 - top_10_acc: 0.9998
Epoch 58/100
80/80 [=====] - 17s 213ms/step - loss: 2.1624 - acc: 0.9852 - top_5_acc: 0.9994 - top_10_acc: 0.9998
Epoch 59/100
80/80 [=====] - 17s 213ms/step - loss: 2.1498 - acc: 0.9869 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 60/100
80/80 [=====] - 17s 213ms/step - loss: 2.1421 - acc: 0.9886 - top_5_acc: 0.9995 - top_10_acc: 0.9999
Epoch 61/100
80/80 [=====] - 17s 212ms/step - loss: 2.1522 - acc: 0.9870 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 62/100
80/80 [=====] - 17s 212ms/step - loss: 2.1498 - acc: 0.9872 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 63/100
80/80 [=====] - 17s 212ms/step - loss: 2.1590 - acc: 0.9853 - top_5_acc: 0.9994 - top_10_acc: 0.9998
Epoch 64/100
80/80 [=====] - 17s 212ms/step - loss: 2.1571 - acc: 0.9859 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 65/100
80/80 [=====] - 17s 213ms/step - loss: 2.1620 - acc: 0.9841 - top_5_acc: 0.9995 - top_10_acc: 0.9999
Epoch 66/100
80/80 [=====] - 17s 212ms/step - loss: 2.1584 - acc: 0.9857 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 67/100
80/80 [=====] - 17s 213ms/step - loss: 2.1479 - acc: 0.9876 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 68/100

80/80 [=====] - 17s 213ms/step - loss: 2.1402 - acc: 0.9890 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 69/100

80/80 [=====] - 17s 213ms/step - loss: 2.1360 - acc: 0.9895 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 70/100

80/80 [=====] - 17s 212ms/step - loss: 2.1389 - acc: 0.9892 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 71/100

80/80 [=====] - 17s 213ms/step - loss: 2.1612 - acc: 0.9847 - top_5_acc: 0.9993 - top_10_acc: 0.9997
Epoch 72/100

80/80 [=====] - 17s 214ms/step - loss: 2.1525 - acc: 0.9873 - top_5_acc: 0.9995 - top_10_acc: 0.9999
Epoch 73/100

80/80 [=====] - 17s 213ms/step - loss: 2.1473 - acc: 0.9869 - top_5_acc: 0.9995 - top_10_acc: 0.9999
Epoch 74/100

80/80 [=====] - 17s 213ms/step - loss: 2.1393 - acc: 0.9890 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 75/100

80/80 [=====] - 17s 213ms/step - loss: 2.1378 - acc: 0.9889 - top_5_acc: 0.9996 - top_10_acc: 0.9998
Epoch 76/100

80/80 [=====] - 17s 214ms/step - loss: 2.1309 - acc: 0.9899 - top_5_acc: 0.9998 - top_10_acc: 0.9999
Epoch 77/100

80/80 [=====] - 17s 213ms/step - loss: 2.1278 - acc: 0.9903 - top_5_acc: 0.9998 - top_10_acc: 1.0000
Epoch 78/100

80/80 [=====] - 17s 212ms/step - loss: 2.1444 - acc: 0.9880 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 79/100

80/80 [=====] - 17s 213ms/step - loss: 2.1399 - acc: 0.9887 - top_5_acc: 0.9997 - top_10_acc: 0.9999
Epoch 80/100

80/80 [=====] - 17s 212ms/step - loss: 2.1372 - acc: 0.9897 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 81/100

80/80 [=====] - 17s 212ms/step - loss: 2.1375 - acc: 0.9893 - top_5_acc: 0.9998 - top_10_acc: 0.9999
Epoch 82/100

80/80 [=====] - 17s 212ms/step - loss: 2.1417 - acc: 0.9879 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 83/100

80/80 [=====] - 17s 213ms/step - loss: 2.1404 - acc: 0.9891 - top_5_acc: 0.9996 - top_10_acc: 0.9998
Epoch 84/100

80/80 [=====] - 17s 212ms/step - loss: 2.1333 - acc: 0.9894 - top_5_acc: 0.9998 - top_10_acc: 0.9999
Epoch 85/100

80/80 [=====] - 17s 212ms/step - loss: 2.1256 - acc: 0.9911 - top_5_acc: 0.9997 - top_10_acc: 1.0000
Epoch 86/100

80/80 [=====] - 17s 212ms/step - loss: 2.1319 - acc: 0.9896 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 87/100

80/80 [=====] - 17s 213ms/step - loss: 2.1290 - acc: 0.9902 - top_5_acc: 0.9998 - top_10_acc: 1.0000
Epoch 88/100

80/80 [=====] - 17s 212ms/step - loss: 2.1353 - acc: 0.9889 - top_5_acc: 0.9997 - top_10_acc: 0.9999
Epoch 89/100

80/80 [=====] - 17s 212ms/step - loss: 2.1451 - acc: 0.9873 - top_5_acc: 0.9995 - top_10_acc: 0.9998
Epoch 90/100

80/80 [=====] - 17s 213ms/step - loss: 2.1329 - acc: 0.9899 - top_5_acc: 0.9997 - top_10_acc: 0.9999
Epoch 91/100

80/80 [=====] - 17s 213ms/step - loss: 2.1321 - acc: 0.9898 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 92/100

80/80 [=====] - 17s 213ms/step - loss: 2.1384 - acc: 0.9879 - top_5_acc: 0.9996 - top_10_acc: 0.9998
Epoch 93/100

80/80 [=====] - 17s 212ms/step - loss: 2.1364 - acc: 0.9889 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 94/100

80/80 [=====] - 17s 212ms/step - loss: 2.1353 - acc: 0.9893 - top_5_acc: 0.9996 - top_10_acc: 0.9999
Epoch 95/100

80/80 [=====] - 17s 212ms/step - loss: 2.1267 - acc: 0.9903 - top_5_acc: 0.9997 - top_10_acc: 0.9999
Epoch 96/100

80/80 [=====] - 17s 213ms/step - loss: 2.1231 - acc: 0.9912 - top_5_acc: 0.9997 - top_10_acc: 0.9998
Epoch 97/100

80/80 [=====] - 17s 213ms/step - loss: 2.1176 - acc: 0.9926 - top_5_acc: 0.9998 - top_10_acc: 0.9999
Epoch 98/100

80/80 [=====] - 17s 214ms/step - loss: 2.1271 - acc: 0.9907 - top_5_acc: 0.9996 - top_10_acc: 0.9998
Epoch 99/100

80/80 [=====] - 17s 213ms/step - loss: 2.1428 - acc: 0.9881 - top_5_acc: 0.9995 - top_10_acc: 0.9999
Epoch 100/100


```
80/80 [=====] - 17s 213ms/step - loss: 2.1362 - acc: 0.9885 - top_5_acc: 0.9996 - top_10_acc: 0.9999
```

```
[ ]: y_val_pred = model.predict({ 'frames': X_val, 'non_empty_frame_idxs':  
    ↪NON_EMPTY_FRAME_IDXS_VAL }, verbose=2).argmax(axis=1)
```

```
446/446 - 6s - 6s/epoch - 12ms/step
```

```
[ ]: y_val_pred.shape
```

```
[ ]: (14248,)
```

```
[ ]: y_val_pred
```

```
[ ]: array([ 47,  30,  54, ..., 145, 142, 238])
```

```
[ ]: y_val
```

```
[ ]: array([ 47,  30,  54, ...,  79, 142, 238], dtype=int32)
```

```
[ ]: import json
```

```
signmap_sub_dir = 'sign_to_prediction_index_map.json'  
signmap_full_file_path = os.path.join(googledrive_dir, signmap_sub_dir)  
# Load the sign to index mapping  
with open(signmap_full_file_path, 'r') as file:  
    sign_to_index = json.load(file)
```

```
[ ]: sign_to_index
```

```
[ ]: {'TV': 0,  
      'after': 1,  
      'airplane': 2,  
      'all': 3,  
      'alligator': 4,  
      'animal': 5,  
      'another': 6,  
      'any': 7,  
      'apple': 8,  
      'arm': 9,  
      'aunt': 10,  
      'awake': 11,  
      'backyard': 12,  
      'bad': 13,  
      'balloon': 14,  
      'bath': 15,  
      'because': 16,  
      'bed': 17,  
      'bedroom': 18,
```

'bee': 19,
'before': 20,
'beside': 21,
'better': 22,
'bird': 23,
'black': 24,
'blow': 25,
'blue': 26,
'boat': 27,
'book': 28,
'boy': 29,
'brother': 30,
'brown': 31,
'bug': 32,
'bye': 33,
'callonphone': 34,
'can': 35,
'car': 36,
'carrot': 37,
'cat': 38,
'cereal': 39,
'chair': 40,
'cheek': 41,
'child': 42,
'chin': 43,
'chocolate': 44,
'clean': 45,
'close': 46,
'closet': 47,
'cloud': 48,
'clown': 49,
'cow': 50,
'cowboy': 51,
'cry': 52,
'cut': 53,
'cute': 54,
'dad': 55,
'dance': 56,
'dirty': 57,
'dog': 58,
'doll': 59,
'donkey': 60,
'down': 61,
'drawer': 62,
'drink': 63,
'drop': 64,
'dry': 65,

'dryer': 66,
'duck': 67,
'ear': 68,
'elephant': 69,
'empty': 70,
'every': 71,
'eye': 72,
'face': 73,
'fall': 74,
'farm': 75,
'fast': 76,
'feet': 77,
'find': 78,
'fine': 79,
'finger': 80,
'finish': 81,
'fireman': 82,
'first': 83,
'fish': 84,
'flag': 85,
'flower': 86,
'food': 87,
'for': 88,
'frenchfries': 89,
'frog': 90,
'garbage': 91,
'gift': 92,
'giraffe': 93,
'girl': 94,
'give': 95,
'glasswindow': 96,
'go': 97,
'goose': 98,
'grandma': 99,
'grandpa': 100,
'grass': 101,
'green': 102,
'gum': 103,
'hair': 104,
'happy': 105,
'hat': 106,
'hate': 107,
'have': 108,
'haveto': 109,
'head': 110,
'hear': 111,
'helicopter': 112,

'hello': 113,
'hen': 114,
'hesheit': 115,
'hide': 116,
'high': 117,
'home': 118,
'horse': 119,
'hot': 120,
'hungry': 121,
'icecream': 122,
'if': 123,
'into': 124,
'jacket': 125,
'jeans': 126,
'jump': 127,
'kiss': 128,
'kitty': 129,
'lamp': 130,
'later': 131,
'like': 132,
'lion': 133,
'lips': 134,
'listen': 135,
'look': 136,
'loud': 137,
'mad': 138,
'make': 139,
'man': 140,
'many': 141,
'milk': 142,
'minemy': 143,
'mitten': 144,
'mom': 145,
'moon': 146,
'morning': 147,
'mouse': 148,
'mouth': 149,
'nap': 150,
'napkin': 151,
'night': 152,
'no': 153,
'noisy': 154,
'nose': 155,
'not': 156,
'now': 157,
'nuts': 158,
'old': 159,

'on': 160,
'open': 161,
'orange': 162,
'outside': 163,
'owie': 164,
'owl': 165,
'pajamas': 166,
'pen': 167,
'pencil': 168,
'penny': 169,
'person': 170,
'pig': 171,
'pizza': 172,
'please': 173,
'police': 174,
'pool': 175,
'potty': 176,
'pretend': 177,
'pretty': 178,
'puppy': 179,
'puzzle': 180,
'quiet': 181,
'radio': 182,
'rain': 183,
'read': 184,
'red': 185,
'refrigerator': 186,
'ride': 187,
'room': 188,
'sad': 189,
'same': 190,
'say': 191,
'scissors': 192,
'see': 193,
'shhh': 194,
'shirt': 195,
'shoe': 196,
'shower': 197,
'sick': 198,
'sleep': 199,
'sleepy': 200,
'smile': 201,
'snack': 202,
'snow': 203,
'stairs': 204,
'stay': 205,
'sticky': 206,

```
'store': 207,  
'story': 208,  
'stuck': 209,  
'sun': 210,  
'table': 211,  
'talk': 212,  
'taste': 213,  
'thankyou': 214,  
'that': 215,  
'there': 216,  
'think': 217,  
'thirsty': 218,  
'tiger': 219,  
'time': 220,  
'tomorrow': 221,  
'tongue': 222,  
'tooth': 223,  
'toothbrush': 224,  
'touch': 225,  
'toy': 226,  
'tree': 227,  
'uncle': 228,  
'underwear': 229,  
'up': 230,  
'vacuum': 231,  
'wait': 232,  
'wake': 233,  
'water': 234,  
'wet': 235,  
'weus': 236,  
'where': 237,  
'white': 238,  
'who': 239,  
'why': 240,  
'will': 241,  
'wolf': 242,  
'yellow': 243,  
'yes': 244,  
'yesterday': 245,  
'yourself': 246,  
'yucky': 247,  
'zebra': 248,  
'zipper': 249}
```

```
[ ]: inverted_mapping = {v: k for k, v in sign_to_index.items()}
```

```
[ ]: # Convert mapping to list
class_names = [inverted_mapping[i] for i in sorted(inverted_mapping)]
```

```
[ ]: from sklearn.metrics import classification_report

print(classification_report(y_val, y_val_pred, target_names=class_names))
```

	precision	recall	f1-score	support
TV	0.76	0.95	0.84	60
after	0.39	0.37	0.38	59
airplane	0.98	1.00	0.99	57
all	0.78	0.64	0.70	59
alligator	0.56	0.77	0.65	61
animal	0.63	0.54	0.58	48
another	0.61	0.69	0.65	51
any	0.56	0.76	0.65	58
apple	0.94	0.89	0.91	53
arm	0.78	0.39	0.52	54
aunt	0.72	0.65	0.69	55
awake	0.44	0.26	0.33	57
backyard	0.76	0.59	0.67	59
bad	0.80	0.67	0.73	60
balloon	0.70	0.79	0.75	63
bath	0.65	0.96	0.78	54
because	0.88	0.83	0.85	59
bed	0.62	0.96	0.76	55
bedroom	0.71	0.61	0.65	61
bee	0.76	0.63	0.69	60
before	0.73	0.61	0.67	54
beside	0.64	0.49	0.55	47
better	0.90	0.77	0.83	57
bird	0.73	0.97	0.83	58
black	0.93	0.88	0.90	58
blow	0.75	0.88	0.81	59
blue	0.83	0.71	0.76	55
boat	0.68	0.81	0.74	53
book	0.71	0.88	0.78	56
boy	0.92	0.91	0.92	54
brother	0.84	0.88	0.86	60
brown	0.84	0.98	0.90	58
bug	0.67	0.90	0.77	62
bye	0.66	0.75	0.70	60
callonphone	1.00	0.72	0.84	57
can	0.52	0.57	0.54	54
car	0.79	0.63	0.70	54
carrot	0.96	0.45	0.62	55
cat	0.54	0.66	0.59	67

cereal	0.79	0.54	0.64	56
chair	0.68	0.85	0.75	52
cheek	0.90	0.74	0.81	61
child	0.76	0.54	0.63	52
chin	0.75	0.78	0.76	58
chocolate	0.83	0.91	0.87	58
clean	0.43	0.87	0.57	53
close	0.76	0.38	0.51	65
closet	0.88	0.57	0.69	61
cloud	0.74	0.75	0.75	61
clown	0.87	0.87	0.87	61
cow	0.92	0.95	0.93	60
cowboy	0.84	0.83	0.83	58
cry	0.87	0.84	0.86	57
cut	0.60	0.79	0.68	57
cute	0.95	0.92	0.93	59
dad	0.81	0.98	0.89	57
dance	0.84	0.70	0.76	46
dirty	0.95	0.75	0.84	53
dog	0.71	0.57	0.63	56
doll	0.98	0.90	0.94	61
donkey	0.88	0.97	0.92	59
down	0.74	0.81	0.78	43
drawer	0.66	0.67	0.67	58
drink	0.90	0.98	0.94	57
drop	0.78	0.43	0.55	49
dry	0.79	0.71	0.75	59
dryer	0.84	0.77	0.80	60
duck	0.64	0.84	0.73	57
ear	0.97	0.48	0.64	60
elephant	0.68	0.58	0.62	59
empty	0.69	0.62	0.65	58
every	0.47	0.48	0.48	50
eye	0.82	0.78	0.80	54
face	0.54	0.72	0.61	53
fall	0.86	0.62	0.72	58
farm	0.87	0.67	0.76	58
fast	0.45	0.27	0.34	51
feet	0.88	0.63	0.74	60
find	0.31	0.56	0.40	59
fine	0.52	0.89	0.66	53
finger	0.72	0.57	0.64	60
finish	0.69	0.78	0.73	60
fireman	0.96	0.83	0.89	60
first	0.85	0.78	0.82	60
fish	0.59	0.71	0.65	56
flag	0.75	0.97	0.84	58
flower	0.93	0.88	0.90	57

food	0.90	0.96	0.93	57
for	0.96	0.95	0.96	57
frenchfries	0.56	0.98	0.71	60
frog	0.98	0.91	0.94	54
garbage	0.97	0.58	0.72	50
gift	0.58	0.84	0.69	58
giraffe	0.91	0.59	0.72	54
girl	0.55	0.42	0.48	52
give	0.59	0.25	0.35	53
glasswindow	0.50	0.35	0.41	54
go	0.86	0.23	0.36	53
goose	0.60	0.05	0.09	62
grandma	0.58	0.75	0.66	57
grandpa	0.98	0.83	0.90	60
grass	0.46	0.47	0.46	60
green	0.79	0.91	0.85	55
gum	0.89	0.93	0.91	59
hair	0.83	0.60	0.70	58
happy	0.82	0.90	0.86	60
hat	0.80	0.81	0.80	43
hate	0.49	0.67	0.57	58
have	0.68	0.88	0.77	49
haveto	0.66	0.85	0.75	55
head	0.79	0.91	0.85	58
hear	0.57	0.56	0.56	68
helicopter	0.66	0.47	0.55	57
hello	0.62	0.79	0.69	57
hen	0.52	0.79	0.63	61
hesheit	0.63	0.77	0.69	60
hide	0.48	0.57	0.52	51
high	0.88	0.93	0.90	54
home	0.98	0.81	0.89	59
horse	0.90	1.00	0.94	60
hot	0.76	0.85	0.80	60
hungry	0.87	0.81	0.84	59
icecream	0.91	0.72	0.80	60
if	0.84	0.93	0.88	55
into	0.67	0.69	0.68	54
jacket	0.70	0.72	0.71	58
jeans	0.76	0.55	0.64	56
jump	0.79	0.63	0.70	59
kiss	0.71	0.63	0.67	59
kitty	0.34	0.26	0.30	65
lamp	0.71	0.59	0.64	61
later	0.59	0.78	0.67	58
like	0.66	0.92	0.77	64
lion	0.90	0.84	0.87	55
lips	0.47	0.58	0.52	60

listen	0.76	0.56	0.65	57
look	0.77	0.44	0.56	62
loud	0.77	0.49	0.60	61
mad	0.63	0.83	0.72	58
make	0.60	0.83	0.70	59
man	0.81	0.78	0.79	59
many	0.40	0.64	0.49	50
milk	0.70	0.78	0.74	58
minemy	0.70	0.97	0.81	60
mitten	0.94	0.48	0.64	62
mom	0.58	0.91	0.71	58
moon	0.93	0.95	0.94	59
morning	0.85	0.93	0.89	55
mouse	0.95	0.92	0.93	60
mouth	0.52	0.22	0.31	60
nap	0.59	0.15	0.24	65
napkin	0.73	0.45	0.56	60
night	0.79	0.58	0.67	45
no	0.68	0.85	0.76	61
noisy	0.92	0.41	0.57	56
nose	0.89	0.88	0.89	58
not	0.59	0.90	0.71	61
now	0.69	0.94	0.80	52
nuts	0.60	0.76	0.67	63
old	0.89	0.63	0.74	54
on	0.78	0.73	0.75	59
open	0.66	0.78	0.71	49
orange	0.94	0.79	0.86	58
outside	0.55	0.48	0.51	61
owie	0.63	0.30	0.40	57
owl	0.96	0.95	0.96	57
pajamas	0.87	0.74	0.80	54
pen	0.69	0.34	0.46	64
pencil	0.61	0.58	0.60	60
penny	0.61	0.46	0.53	54
person	0.36	0.40	0.38	35
pig	0.89	0.92	0.91	53
pizza	0.58	0.38	0.46	66
please	0.90	0.78	0.84	60
police	0.98	0.87	0.92	60
pool	0.90	0.48	0.63	56
potty	0.81	0.94	0.87	53
pretend	0.82	0.84	0.83	61
pretty	0.56	0.52	0.54	58
puppy	0.80	0.52	0.63	62
puzzle	0.68	0.78	0.73	50
quiet	0.70	0.66	0.68	61
radio	0.89	0.65	0.75	60

rain	0.87	0.65	0.74	60
read	0.82	0.75	0.78	53
red	0.90	0.74	0.81	58
refrigerator	0.88	0.53	0.66	55
ride	0.90	0.34	0.49	53
room	0.44	0.67	0.53	60
sad	0.77	0.92	0.84	60
same	0.81	0.80	0.80	59
say	0.38	0.64	0.48	55
scissors	0.60	0.73	0.66	60
see	0.62	0.98	0.76	60
shhh	0.84	0.95	0.89	60
shirt	0.88	0.93	0.90	54
shoe	0.92	0.90	0.91	60
shower	0.72	0.62	0.67	55
sick	0.87	0.70	0.77	56
sleep	0.37	0.67	0.48	57
sleepy	0.54	0.71	0.61	63
smile	0.73	0.73	0.73	56
snack	0.81	0.60	0.69	58
snow	0.94	0.84	0.89	58
stairs	0.83	0.63	0.72	60
stay	0.69	0.75	0.72	56
sticky	0.40	0.79	0.53	53
store	0.82	0.93	0.87	55
story	0.94	0.56	0.70	59
stuck	0.90	0.95	0.92	56
sun	0.62	0.50	0.55	58
table	0.71	0.92	0.80	51
talk	0.97	0.50	0.66	58
taste	0.73	0.91	0.81	57
thankyou	0.38	0.90	0.53	52
that	0.58	0.60	0.59	55
there	0.52	0.46	0.49	50
think	0.58	0.76	0.66	55
thirsty	0.93	0.86	0.89	49
tiger	0.89	0.91	0.90	56
time	0.82	0.84	0.83	50
tomorrow	0.87	0.49	0.63	55
tongue	0.49	0.81	0.61	58
tooth	0.51	0.62	0.56	53
toothbrush	0.67	0.67	0.67	57
touch	0.68	0.60	0.64	57
toy	0.98	0.80	0.88	55
tree	0.81	0.84	0.82	56
uncle	0.93	0.87	0.90	60
underwear	0.86	0.62	0.72	61
up	0.55	0.87	0.68	60

vacuum	0.45	0.48	0.47	52
wait	0.51	0.69	0.59	51
wake	0.41	0.64	0.50	58
water	0.81	0.95	0.87	62
wet	0.46	0.45	0.46	66
weus	0.69	0.88	0.78	57
where	0.80	0.93	0.86	56
white	0.88	0.73	0.80	60
who	0.86	0.80	0.83	60
why	0.82	0.71	0.76	59
will	0.87	0.59	0.70	56
wolf	0.71	0.61	0.66	57
yellow	0.80	0.83	0.81	58
yes	0.63	0.74	0.68	62
yesterday	0.78	0.65	0.71	60
yourself	0.88	0.75	0.81	57
yucky	0.44	0.34	0.38	56
zebra	0.73	0.87	0.79	61
zipper	0.75	0.66	0.70	32
accuracy			0.71	14248
macro avg	0.74	0.71	0.71	14248
weighted avg	0.74	0.71	0.71	14248

The overall ASL Transformer designed by Wijkhuizen, M. are shown in APPENDIX 1. After training and valuation, the Model performance is shown in APPENDIX 2. Wijkhuizen, M.'s transformer model has an overall 0.71 F1 score with a weighted precision of 0.74 and a weighted recall of 0.71. It has outperformed any other model types that we tried. Some ASL word predictions perform better than others; for example, airplane, apple, owl etc., have F1 scores higher than 0.90. and other words like kitty, yucky, and suffer under the F1 score lower than 0.40. However, most words' F1 scores are higher than 0.60, so we could use this model for a real-life application with some limitations.

Reference:

Wijkhuizen, M. (2023, April 04). GISLR TF Data Processing & Transformer Training. Kaggle. <https://www.kaggle.com/code/markwijkhuizen/gislr-tf-data-processing-transformer-training>

Kumar, A. (2023, February 1). The transformer model and its applications: Understanding attention mechanism. Medium. Retrieved from <https://medium.com/@mittal.atul06/the-transformer-model-and-its-applications-understanding-attention-mechanism-c37e6e3a76dc>

Huang, Z., Liang, D., Xu, P., & Xiang, B. (2020). Improve Transformer Models with Better Relative Position Embeddings. arXiv preprint arXiv:2009.13658